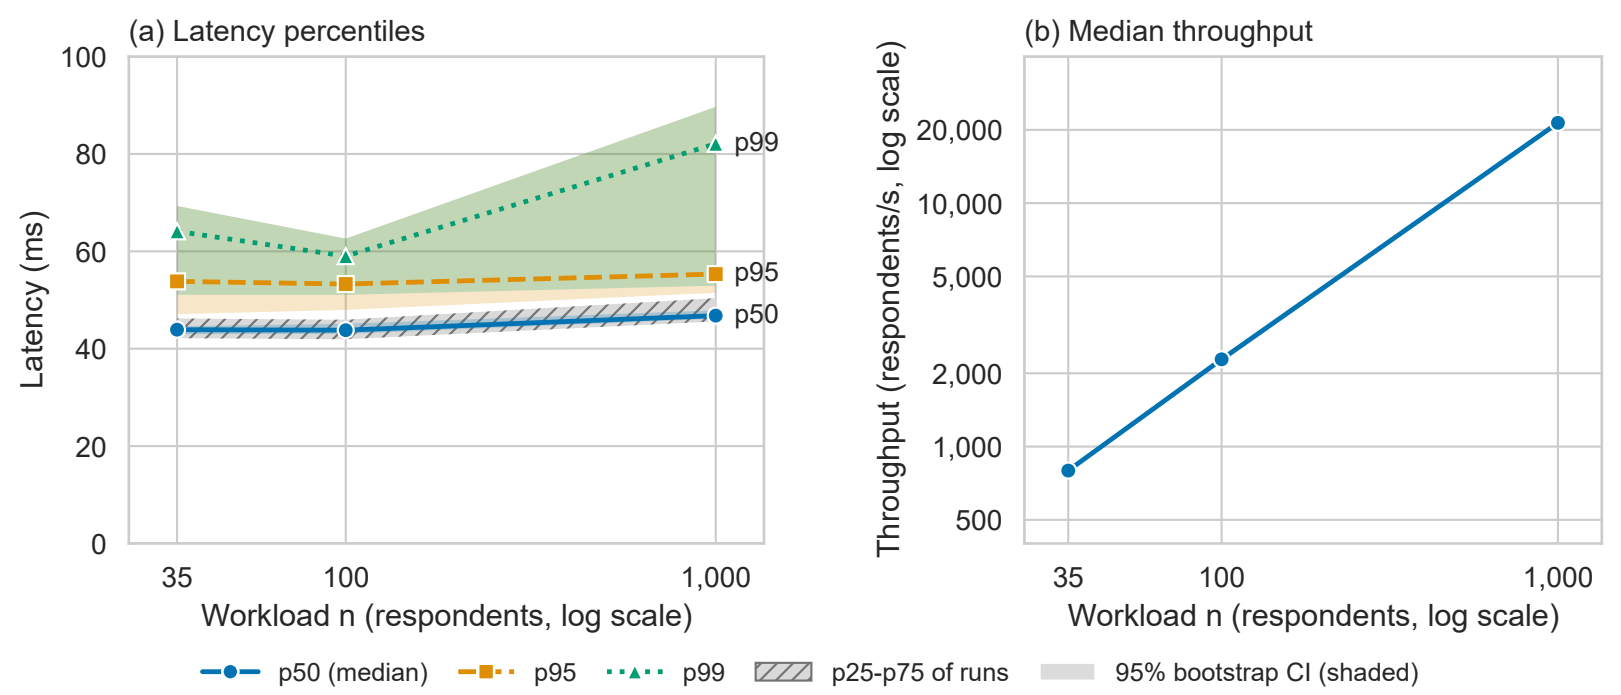
(a) Latency percentiles
Latency (ms)
100
80
60
40
20
0
p99
p95
p50
35
100
1,000
Workload n (respondents, log scale)
(b) Median throughput
Throughput (respondents/s, log scale)
20,000
10,000
5,000
2,000
1,000
500
35
100
1,000
Workload n (respondents, log scale)
p50 (median)
p95
p99
p25-p75 of runs
95% bootstrap CI (shaded)